

Private RAG AWS GPU Deployment - Live-Scripting Workflow

Andreas Wittmann

Andreas Wittmann IT-Beratung GmbH

andreas.wittmann@anwi.gmbh

2025-11-18

Contents

1	Summary	1
2	Building Privacy-Focused RAG Platform with GPU Acceleration	1
2.1	Project Overview	1
2.2	Live-Scripting Methodology	1
2.2.1	Implementation Options	2
2.3	Project Components	2
2.4	Expected Deployment Outcome	2
3	Foundation & Prerequisites	2
3.1	Environment Setup	2
3.1.1	Setting the Project Foundation	2
3.1.2	Validating AWS Prerequisites	3
3.2	Summary	3
4	Infrastructure as Code	3
4.1	SSH Key Generation for Secure Access	3
4.2	Terraform Configuration and Planning	4
4.3	Infrastructure Deployment Planning	5
4.4	Infrastructure Deployment Execution	5
4.5	SSH Configuration for Easy Access	6
4.6	GPU Instance Selection and Upgrading	6
4.6.1	Recommended GPU EC2 Instance Types for RAG Workloads	6
4.6.2	Upgrading to a Different GPU Instance Type with Terraform	7
4.7	Infrastructure Live-Cycle Operations	8
4.7.1	Recommended GPU EC2 Instance Types for RAG Workloads	9
4.8	Summary	9
5	GPU Verification	9
6	Development Environment Setup	10
6.1	Complete Docker Engine and Tools Installation	10
6.1.1	Docker Engine Installation	10
6.1.2	Docker Service Configuration	11
6.1.3	Lazydocker Installation	11

6.1.4	GitHub CLI Installation	12
6.1.5	Htop Installation	13
6.1.6	Docker and Docker Compose Verification	13
6.1.7	Python Development Tools Installation	13
6.1.8	Development Environment Summary	14
6.2	Summary	15
7	Core RagFlow Deployment	15
7.1	Basic Ragflow Setup	15
7.2	HTTPS Setup with Self-Signed Certificate	16
7.2.1	Step 1: Generate Self-Signed SSL Certificate**	16
7.2.2	Step 2: Modify Nginx Configuration**	17
7.2.3	Step 3: Update Docker Compose Configuration**	17
7.2.4	Step 4: Restart RAGFlow Services**	18
7.2.5	Step 5: Test HTTPS Access**	18
7.3	Use Case: Access RagFlow via SSH-Tunnel	19
7.4	Summary	19
8	Advanced Features - Document Processing	19
8.1	Document Ingestion and Processing	19
8.1.1	Upload Documents	19
8.1.2	Configure Knowledge Base	19
8.1.3	Test RAG Queries	19
8.2	Summary	20
9	Infrastructure Cleanup	20
9.1	Cleanup Philosophy	20
9.2	Stop RagFlow Services	20
9.3	Terminate GPU EC2 Instance with Terraform	20
9.4	Clean Up Local SSH Configuration	21
9.5	Clean Up Terraform State Files	21
9.6	Verify Complete Cleanup	22
9.7	Cost Optimization Notes	23
9.8	Summary	23

1 Summary

This document outlines a live-scripting methodology for deploying RagFlow, an open-source RAG platform, on Ubuntu 24.04 with GPU acceleration, through executable code blocks that ensure reproducible infrastructure setup. It covers Terraform infrastructure provisioning with pre-configured Deep Learning AMI, Docker containerization, and RagFlow deployment with GPU support.

- **HTML Version:** [private-rag-aws-gpu.html](#)
- **PDF Version:** [private-rag-aws-gpu.pdf](#)
- **GitHub:** <https://github.com/andreaswittmann/private-rag/tree/main>

2 Building Privacy-Focused RAG Platform with GPU Acceleration

2.1 Project Overview

This project uses the live-scripting methodology to create an example configuration of the open-source project RagFlow on Ubuntu 24.04 with GPU support. Live-scripting represents an intermediate approach between manual command-line work and fully automated CI/CD pipelines, enabling the development of documented, shareable, and reproducible solutions.

Using Infrastructure as Code methods with Terraform, an AWS environment with GPU-enabled EC2 instance running Ubuntu 24.04 is established as a platform for RagFlow installation. The Deep Learning AMI comes pre-configured with NVIDIA drivers and CUDA toolkit, eliminating manual GPU setup. The example demonstrates all necessary configuration and software tool installations required for the project.

2.2 Live-Scripting Methodology

The deployment uses live-scripting methodology to transform static documentation into executable workflows. Code blocks execute directly in Emacs (F4 key) or copy to any terminal, creating a bridge between documentation and execution.

Based on the [live-scripting methodology](<https://github.com/andreaswittmann/live-scripting>), this approach provides:

- Executable documentation with current and functional code blocks
- Multi-format publishing from single source to Org, HTML, and PDF
- Reproducible deployments with consistent infrastructure setup
- Version control integration through plain-text format

2.2.1 Implementation Options

Emacs Users Configure live-scripting with this Emacs Lisp function:

```
(defun send-line-to-vterm ()
  "Send region if active, or current line to vterm buffer."
  (interactive)
  (if (region-active-p)
      (send-region "*vterm*" (region-beginning) (region-end))
      (my-select-current-line)
      (send-region "*vterm*" (region-beginning) (region-end)))
  (deactivate-mark))

(global-set-key [f4] 'send-line-to-vterm)
```

Non-Emacs Users Copy code blocks to any terminal - the methodology maintains precision across tools.

The deployment follows phases from environment setup through GPU configuration and RagFlow deployment, with each step documented using executable code blocks for reproducibility.

2.3 Project Components

- Self-hosted RagFlow with local document processing for privacy control
- GPU-accelerated infrastructure with Ubuntu 24.04 and pre-installed NVIDIA drivers/CUDA
- Production-ready foundation with SSL/HTTPS, Docker containerization, and infrastructure automation

- Advanced RAG capabilities including document ingestion, vector search, and knowledge base management
- Reproducible methodology using live-scripting workflow for consistent deployment

2.4 Expected Deployment Outcome

The workflow produces a functioning HTTPS-enabled RagFlow instance on AWS with GPU acceleration, capable of document processing and RAG-based AI responses with full user control and enhanced performance.

3 Foundation & Prerequisites

3.1 Environment Setup

These steps establish the development environment and validate prerequisites.

3.1.1 Setting the Project Foundation

This establishes the base project directory and validates AWS network prerequisites.

```
# Listing: Project environment setup
# Establishing the base directory for the RAG project

bash
PROJECT_DIR="/Users/${whoami}/LocalProjects/RagFlow/private-rag"
cd "$PROJECT_DIR"
pwd
ls -la

# Verify we're in the right location
echo "Working in: $PROJECT_DIR"

# Set AWS profile for this session
export AWS_PROFILE=lab-a-north
echo $AWS_PROFILE

## create default VPC if it does not exist
aws ec2 create-default-vpc
## check if the default vpc is created
aws ec2 describe-vpcs --filters "Name=isDefault,Values=true" --query "Vpcs[0].VpcId" --output
→ text
```

Result: Project environment is configured and AWS default VPC is available for deployment.

3.1.2 Validating AWS Prerequisites

This confirms AWS environment configuration for deployment requirements.

```
# Listing: AWS environment validation
# Ensuring AWS access and basic services availability

# Test basic AWS connectivity
```

```
aws s3 ls

# Verify AWS identity and permissions
aws sts get-caller-identity

# Verify SSH key is available for EC2 access
ls -la ~/.ssh/

# Validate required tools are installed
terraform --version
git --version
docker --version
```

Result: AWS credentials and required tools are validated for infrastructure deployment.

3.2 Summary

The foundation phase establishes the development environment and validates prerequisites. Project directory structure is configured, AWS credentials verified, and required tools confirmed operational. This groundwork enables reproducible infrastructure deployment through live-scripting methodology.

4 Infrastructure as Code

4.1 SSH Key Generation for Secure Access

This creates dedicated SSH keypairs for EC2 instance access and stores the public key for Terraform.

```
# Listing: SSH key generation for EC2 access
# Creating dedicated SSH keys for secure instance access

cd "$PROJECT_DIR"
pwd;
# Generate a new SSH key pair specifically for RagFlow GPU
ssh-keygen -t rsa -b 4096 -f ~/.ssh/ragflow-gpu_key -N ""
y
# Display the public key for use in Terraform variables
cat ~/.ssh/ragflow-gpu_key.pub

# Store public key in environment variable
export TF_VAR_SSH_PUBLIC_KEY=$(cat ~/.ssh/ragflow-gpu_key.pub)
echo "SSH public key configured: $TF_VAR_SSH_PUBLIC_KEY"
```

Result: SSH keypair is generated and public key is exported for infrastructure provisioning.

4.2 Terraform Configuration and Planning

This prepares the Terraform configuration files and validates the infrastructure setup for GPU-enabled instance.

```
# Listing: Terraform environment preparation
# Setting up infrastructure configuration for GPU instance
```

```
cd "$PROJECT_DIR/terraform-gpu"
pwd
# Create terraform variables file with GPU configuration
cat > terraform.tfvars << EOF
# Generated tfvars file - $(date)
aws_region      = "eu-north-1"
instance_type   = "g4dn.xlarge"
root_volume_size = 100
root_volume_type = "gp3"
allowed_ip      = "0.0.0.0/0"
ssh_public_key  = "${TF_VAR_SSH_PUBLIC_KEY}"
ec2_name        = "EC2-RagFlow-GPU"
environment     = "development"
project         = "ragflow-gpu"
ssh_cidr_blocks = ["0.0.0.0/0"]
EOF

# Review our configuration
cat terraform.tfvars

# Initialize Terraform
terraform init

# Format and validate our configuration
terraform fmt
terraform validate
```

Result: Terraform configuration is initialized and validated for GPU-enabled infrastructure deployment.

4.3 Infrastructure Deployment Planning

This creates the deployment plan and reviews resource changes before applying them.

```
# Listing: Terraform deployment planning
# Creating and reviewing the GPU infrastructure deployment plan

# Generate deployment plan
terraform plan -out=tfplan -var-file=terraform.tfvars

# Create human-readable plan
terraform show -json tfplan > tfplan.json
cat tfplan.json | jq > tfplan.pretty.json

# Review what resources will be created
echo "Resources to be created:"
cat tfplan.json | jq '.resource_changes[] | {address: .address, action: .change.actions[0]}'

# Summary of planned changes
cat tfplan.json | jq
→ '.resource_changes | group_by(.change.actions[0]) | map({action: .[0].change.actions[0], count: l
```

Result: Deployment plan is generated and resource changes are reviewed for approval.

4.4 Infrastructure Deployment Execution

This executes the Terraform plan to create the AWS infrastructure resources with GPU support.

```
# Listing: AWS GPU infrastructure deployment
# Creating the GPU-enabled infrastructure resources

# Execute the deployment plan
terraform apply tfplan
terraform apply
yes

# Capture and display outputs
terraform output

# List our EC2 instances
aws ec2 describe-instances \
--query
'Reservations[*].Instances[*].[InstanceId, InstanceType, Tags[?Key==`Name`]![0].Value, State.Name]'
--output text

# Get our instance details - get only the first/latest instance ID
instance_id=$(aws ec2 describe-instances \
--filters "Name=tag:Name,Values=EC2-RagFlow-GPU"
--query "Name=instance-state-name,Values=running,pending,stopping,stopped" \
--query 'Reservations[*].Instances[*].[InstanceId]' \
--output text | head -n1)
echo "Instance ID: $instance_id"

public_ip=$(aws ec2 describe-instances \
--instance-ids "$instance_id" \
--query 'Reservations[*].Instances[*].[PublicIpAddress]' \
--output text)
echo "Public IP: $public_ip"
exit
```

Result: AWS GPU-enabled infrastructure is deployed and EC2 instance details are captured for subsequent configuration.

4.5 SSH Configuration for Easy Access

This establishes convenient SSH access configuration and tests connectivity to the deployed GPU instance.

```
# Listing: SSH configuration setup for GPU instance
# Establishing convenient SSH access

pwd
# Test initial SSH connection
ssh -i ~/.ssh/ragflow-gpu_key ubuntu@$public_ip 'whoami && pwd'
```

```

yes
# Create SSH config entry for easy access
open ~/.ssh/config # we may delete an old entry first, manually
cat >> ~/.ssh/config << EOF
# SSH over EC2 - RagFlow GPU Project
Host EC2-RagFlow-GPU
    HostName $public_ip
    User ubuntu
    IdentityFile ~/.ssh/ragflow-gpu_key
EOF
## use a perl one-liner to grep these lines from the ssh-config file.
perl -nE 'print if /Host EC2-RagFlow-GPU/ .. /yesEOF/' ~/.ssh/config
#bbedit ~/.ssh/config

# Test SSH with hostname
ssh EC2-RagFlow-GPU 'whoami && date'

```

Result: SSH configuration is established and remote access to the GPU EC2 instance is verified.

4.6 GPU Instance Selection and Upgrading

4.6.1 Recommended GPU EC2 Instance Types for RAG Workloads

Here is a list of recommended GPU-enabled EC2 instance types suitable for RAG workloads with RagFlow, focusing on instances with NVIDIA GPUs optimized for AI/ML tasks. All recommendations are for the `eu-north-1` region, with approximate on-demand pricing (as of late 2024; prices may vary—check AWS Pricing Calculator for current rates).

Instance Type	GPU	VRAM (GB)	vCPUs	RAM (GB)	Network Speed	Approx. Price/Hour (USD)
g4dn.xlarge	1x T4	16	4	16	Up to 25 Gbps	\$0.71
g4dn.2xlarge	1x T4	16	8	32	Up to 25 Gbps	\$1.05
g5.xlarge	1x A10G	24	4	16	Up to 10 Gbps	\$1.21
g6e.xlarge	1x L40S	48	4	32	Up to 20 Gbps	\$1.22

Recommendations

- **For most RAG setups:** Start with `g4dn.xlarge` for a balance of GPU performance, memory, and cost.
- **For high-performance needs:** Opt for `g5.xlarge` or `g6e.xlarge` if you need more VRAM for larger models.
- **Considerations:** GPU instances are more expensive but provide significant performance improvements for embeddings and inference tasks.

4.6.2 Upgrading to a Different GPU Instance Type with Terraform

The best way to change your EC2 instance type is to update your Terraform configuration and apply the changes. This approach is automated, ensures consistency, and avoids manual errors.

Steps to Change Instance Type

1. Update the Terraform Variables:

- Edit `terraform-gpu/terraform.tfvars` and change `instance_type = "g4dn.xlarge"` to your desired instance type (e.g., `instance_type = "g5.xlarge"`).

- Alternatively, update the default in `terraform-gpu/variables.tf` if you prefer not to modify `tfvars`.

2. Apply the Changes:

- Navigate to the `terraform-gpu/` directory in your terminal.
- Run `terraform plan` to preview the changes (it will show the instance type update).
- Run `terraform apply` to execute the upgrade. Terraform will:
 - Stop the existing instance.
 - Change the instance type.
 - Restart the instance with the new configuration.
- This process typically takes 5-10 minutes, depending on AWS availability.

Why This Method Is Best

- **Automation:** Terraform handles the entire process, ensuring consistency and avoiding manual errors.
- **Minimal Downtime:** The stop/change/start cycle is efficient for instance type changes.
- **No Data Loss:** Your EBS root volume and Elastic IP remain attached.
- **Cost:** Monitor via AWS Cost Explorer. GPU instances have different pricing.
- **Alternatives Considered:**
 - Manual AWS Console change: Faster for one-off updates but not recommended for Terraform-managed infrastructure, as it can cause drift.
 - Creating a new instance: Slower and requires data migration.

```
# Listing: Changing GPU instance type with Terraform
# Updating instance configuration for different GPU requirements

# Navigate to the Terraform GPU directory
cd "$PROJECT_DIR/terraform-gpu"

# Update the instance type in terraform.tfvars (replace 'g5.xlarge' with your desired type)
sed -i 's/instance_type      = "g4dn.xlarge"/instance_type      = "g5.xlarge"/g' terraform.tfvars

# Review the updated configuration
cat terraform.tfvars

# Generate and review the deployment plan
terraform plan -out=tfplan -var-file=terraform.tfvars

# Apply the changes to upgrade the instance
terraform apply tfplan

# Verify the instance has been upgraded
aws ec2 describe-instances \
  --filters "Name=tag:Name,Values=EC2-RagFlow-GPU" \
  --query 'Reservations[*].Instances[*].[InstanceType, State.Name]' \
  --output text
```

Result: GPU instance type has been successfully changed to the new configuration with updated GPU capabilities.

4.7 Infrastructure Live-Cycle Operations

This section provides procedures for managing the GPU EC2 instance lifecycle - starting and stopping the instance when returning to work sessions after periods of inactivity.

```
exit

bash
pwd

# Initialize project environment
PROJECT_DIR="/Users/${whoami}/LocalProjects/RagFlow/private-rag"
cd "$PROJECT_DIR"
export AWS_PROFILE=lab-a-north

# Get instance ID
instance_id=$(aws ec2 describe-instances --filters "Name>tag:Name,Values=EC2-RagFlow-GPU"
→ --query 'Reservations[*].Instances[*].[InstanceId]' --output text | head -n1)
echo $instance_id

### STATUS: Get current instance state
aws ec2 describe-instances --instance-ids "$instance_id" --query
→ 'Reservations[*].Instances[*].[State.Name]' --output text

### START: Start and wait for running state
aws ec2 start-instances --instance-ids "$instance_id"
while [ "$(aws ec2 describe-instances --instance-ids "$instance_id" --query
→ "Reservations[0].Instances[0].State.Name" --output text)" != "running" ]; do echo -n ".";
→ sleep 1; done; echo " running"

# Get public IP
public_ip=$(aws ec2 describe-instances --instance-ids "$instance_id" --query
→ 'Reservations[*].Instances[*].[PublicIpAddress]' --output text)
echo "Instance running at: $public_ip"

# Test SSH connection
ssh EC2-RagFlow-GPU 'whoami && date && uptime'

### STOP: stop the instance and wait for stopped state
aws ec2 stop-instances --instance-ids "$instance_id"
while [ "$(aws ec2 describe-instances --instance-ids "$instance_id" --query
→ "Reservations[0].Instances[0].State.Name" --output text)" != "stopped" ]; do echo -n ".";
→ sleep 1; done; echo " stopped"
```

Result: GPU EC2 instance is started and ready for use. Stop commands are provided as well.

4.7.1 Recommended GPU EC2 Instance Types for RAG Workloads

Here is a list of recommended GPU-enabled EC2 instance types suitable for RAG workloads with RagFlow, focusing on instances with NVIDIA GPUs optimized for AI/ML tasks. All recommendations are for the `eu-central-1` region, with approximate on-demand pricing (as of late 2024; prices may vary—check AWS Pricing Calculator for current rates).

Instance Type	GPU	VRAM (GB)	vCPUs	RAM (GB)	Network Speed	Approx. Price/Hour (USD)
g4dn.xlarge	1x T4	16	4	16	Up to 25 Gbps	\$0.71
g4dn.2xlarge	1x T4	16	8	32	Up to 25 Gbps	\$1.05
g5.xlarge	1x A10G	24	4	16	Up to 10 Gbps	\$1.21
g6e.xlarge	1x L40S	48	4	32	Up to 20 Gbps	\$1.22

Recommendations

- **For most RAG setups:** Start with `g4dn.xlarge` for a balance of GPU performance, memory, and cost.
- **For high-performance needs:** Opt for `g5.xlarge` or `g6e.xlarge` if you need more VRAM for larger models.
- **Considerations:** GPU instances are more expensive but provide significant performance improvements for embeddings and inference tasks.

4.8 Summary

Infrastructure automation through Terraform provisions the AWS environment with GPU support. The EC2 instance operates with proper network security configuration and NVIDIA GPU capabilities. SSH access is established through dedicated keypairs. The infrastructure exists as version-controlled code, enabling consistent reproduction across deployments.

5 GPU Verification

This verifies that the pre-installed GPU drivers, CUDA, and Docker GPU support are working correctly.

```
# Listing: GPU verification
# Confirming pre-installed GPU components are operational

ssh EC2-RagFlow-GPU

# Check NVIDIA driver and GPU status
nvidia-smi

# Check CUDA version
nvcc --version

# Test GPU in Docker
docker run --rm --gpus all nvidia/cuda:12.9.0-base-ubuntu24.04 nvidia-smi

exit
```

Result: Pre-installed GPU drivers, CUDA toolkit, and Docker GPU support are verified as fully operational.

6 Development Environment Setup

6.1 Complete Docker Engine and Tools Installation

This establishes the complete development environment on the Ubuntu 24.04 GPU instance, including Docker Engine, container management tools, GitHub CLI, and Python development utilities.

6.1.1 Docker Engine Installation

Install Docker Engine with proper repository configuration and user permissions.

```
# Listing: Docker Engine installation on GPU instance
# Setting up containerization platform for GPU-accelerated applications

ssh EC2-RagFlow-GPU
bash

# Remove any old Docker versions that may exist
sudo apt remove docker docker-engine docker.io containerd runc 2>/dev/null || true

# Update system packages
sudo apt update -y
sudo apt install -y ca-certificates curl gnupg lsb-release

# Add Docker's official GPG key
sudo install -m 0755 -d /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o
→ /etc/apt/keyrings/docker.gpg
sudo chmod a+r /etc/apt/keyrings/docker.gpg

# Add Docker repository
echo \
"deb [arch=$(dpkg
→ --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] https://download.docker.com/linux/u
$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null

# Install Docker packages
sudo apt update
sudo apt install -y docker-ce docker-ce-cli containerd.io docker-buildx-plugin
→ docker-compose-plugin

# Configure Docker to run without sudo (important for CI/CD)
sudo usermod -aG docker $USER

# Apply group changes without logout
newgrp docker

# Test Docker installation
docker run hello-world

# Verify Docker Compose
docker compose version

exit
```

Result: Docker Engine is installed with proper GPG key verification and user permissions configured for passwordless operation.

6.1.2 Docker Service Configuration

Enable Docker as a system service for automatic startup and management.

```
# Listing: Docker service setup
# Configuring Docker for automatic startup and monitoring

ssh EC2-RagFlow-GPU

# Enable Docker service to start on boot
sudo systemctl enable docker

# Start Docker service
sudo systemctl start docker

# Verify service status
sudo systemctl status docker --no-pager

exit
```

Result: Docker service is configured to start automatically on EC2 instance boot.

6.1.3 Lazydocker Installation

Install Lazydocker for convenient terminal-based Docker container management.

```
# Listing: Lazydocker installation
# Installing terminal UI for Docker management

ssh EC2-RagFlow-GPU

# Download and install Lazydocker
curl
→ https://raw.githubusercontent.com/jesseduffield/lazydocker/master/scripts/install_update_linux.sh
→ | bash

# Add Lazydocker to PATH
echo 'export PATH="$HOME/.local/bin:$PATH"' >> ~/.bashrc

# Reload bashrc to apply PATH changes
source ~/.bashrc

# Verify installation
lazydocker --version

exit
```

Result: Lazydocker is installed and available in the PATH for terminal-based Docker management.

6.1.4 GitHub CLI Installation

Install the GitHub CLI tool for repository and runner management.

```
# Listing: GitHub CLI installation
# Installing gh command-line tool for GitHub integration
```

```
ssh EC2-RagFlow-GPU

# Update package list
sudo apt update

# Install dependencies
sudo apt install -y curl unzip

# Add GitHub CLI GPG key
curl -fsSL https://cli.github.com/packages/githubcli-archive-keyring.gpg | sudo dd
→ of=/usr/share/keyrings/githubcli-archive-keyring.gpg
sudo chmod go+r /usr/share/keyrings/githubcli-archive-keyring.gpg

# Add GitHub CLI repository
echo "deb [arch=$(dpkg
→ --print-architecture) signed-by=/usr/share/keyrings/githubcli-archive-keyring.gpg] https://cli.gi
→ | sudo tee /etc/apt/sources.list.d/github-cli.list > /dev/null

# Install GitHub CLI
sudo apt update
sudo apt install -y gh

# Verify installation
gh --version

exit
```

Result: GitHub CLI is installed and ready for repository and runner management operations.

6.1.5 Htop Installation

Install htop for interactive process monitoring and system resource visualization.

```
# Listing: Htop installation
# Installing interactive system monitor

ssh EC2-RagFlow-GPU

# Install htop
sudo apt install -y htop

# Verify installation
htop --version

exit
```

Result: Htop is installed and available for system monitoring and process management.

6.1.6 Docker and Docker Compose Verification

Verify all Docker components are properly installed and functional.

```
# Listing: Docker installation verification
# Confirming all Docker components are operational

ssh EC2-RagFlow-GPU

# Check Docker version
docker --version

# Check Docker Compose version (new recommended format)
docker compose version

# Verify Docker daemon is running
docker ps

# List available Docker images
docker images

exit
```

Result: Docker Engine and Docker Compose are verified as fully operational.

6.1.7 Python Development Tools Installation

Install uv (fast Python package installer) and development libraries required for Python projects.

```
# Listing: Python development tools installation
# Installing uv and development dependencies

ssh EC2-RagFlow-GPU

# Install uv (fast Python package installer and runner)
curl -LsSf https://astral.sh/uv/install.sh | sh

# Verify uv installation
uv --version
uvx --version

# Install ICU development libraries and build tools
sudo apt update
sudo apt install -y \
    libicu-dev \
    pkg-config \
    build-essential

# Verify build tools
gcc --version
pkg-config --version

exit
```

Result: Python development tools including uv, ICU libraries, and build essentials are installed and verified.

6.1.8 Development Environment Summary

Verify the complete development environment is ready for deployment.

```
# Listing: Development environment verification
# Confirming all tools are installed and operational

ssh EC2-RagFlow-GPU

# Display installed versions
echo "==== Development Environment Status ===="
echo "Docker:"
docker --version
echo ""
echo "Docker Compose:"
docker compose version
echo ""
echo "GitHub CLI:"
gh --version
echo ""
echo "Python Tools:"
uv --version
echo ""
echo "Build Tools:"
gcc --version | head -1
echo ""
echo "CUDA:"
/usr/local/cuda/bin/nvcc --version | head -1
echo ""
echo "GPU:"
nvidia-smi --query-gpu=name --format=csv,noheader
echo ""
echo "System Information:"
uname -a
echo ""
echo "Available Disk Space:"
df -h /
echo ""
echo "GPU Memory:"
nvidia-smi --query-gpu=memory.total --format=csv,noheader

exit
```

Result: Complete development environment with GPU support is verified and ready for RagFlow deployment and CI/CD operations.

6.2 Summary

The development environment setup establishes necessary tools for containerized application deployment and CI/CD workflows on GPU-enabled infrastructure. Docker Engine provides the containerization platform, Lazydock offers terminal management, GitHub CLI enables repository integration, and Python development tools support application requirements. CUDA and NVIDIA drivers enable GPU acceleration for enhanced performance.

7 Core RagFlow Deployment

7.1 Basic Ragflow Setup

This section deploys RagFlow on the GPU-enabled EC2 instance using Docker Compose following the QuickStart guide.

```
# Listing: RagFlow deployment on GPU instance
# Installing and starting RagFlow RAG platform with GPU support

# SSH into the GPU EC2 instance
ssh EC2-RagFlow-GPU

# Clone the RagFlow repository
cd ~; pwd
git clone https://github.com/infiniflow/ragflow.git

# Navigate to the RagFlow directory
cd ~/ragflow/docker

# Start RagFlow using Docker Compose
docker compose up -d

# Verify the deployment
docker ps

# Getting SERVICE_IP
public_ip=$(curl -s http://checkip.amazonaws.com)
echo $public_ip

# RagFlow URL
echo "RagFlow is accessible at: http://$public_ip:80"

# Shutting down RagFlow (optional)
docker compose stop
# Clean up RagFlow (optional)
docker compose down -v # removes volumes as well
cd ~; rm -rf ragflow

exit
```

Result: RagFlow is successfully deployed and running on the GPU-enabled EC2 instance.

7.2 HTTPS Setup with Self-Signed Certificate

This section provides step-by-step instructions for setting up HTTPS access to RAGFlow using a self-signed SSL certificate on the GPU instance.

Prerequisites

- Docker and Docker Compose installed on your GPU server
- Ports 80 and 443 open on your server
- Server's public IP address (replace <SERVERIP> with your actual IP throughout these instructions)

- Basic knowledge of command-line operations

7.2.1 Step 1: Generate Self-Signed SSL Certificate**

Generate a self-signed certificate for your server's IP address.

```
bash
pwd
ssh EC2-RagFlow-GPU
cd
cd ~/ragflow/docker/nginx

# Retrieve the SERVER_IP
curl -s http://checkip.amazonaws.com
export SERVER_IP=$(curl -s http://checkip.amazonaws.com)
echo $SERVER_IP

# Create a directory for SSL certificates
mkdir -p ssl

# Generate a private key
openssl genrsa -out ssl/privkey.pem 2048

# Generate a certificate signing request (CSR) for your IP
openssl req -new -key ssl/privkey.pem -out ssl/cert.csr -subj
↪ "/C=US/ST=State/L=City/O=Organization/CN=localhost"
openssl req -new -key ssl/privkey.pem -out ssl/cert.csr -subj
↪ "/C=US/ST=State/L=City/O=Organization/CN=$SERVER_IP"

# Generate the self-signed certificate valid for 365 days
openssl x509 -req -days 365 -in ssl/cert.csr -signkey ssl/privkey.pem -out ssl/fullchain.pem

# Set appropriate permissions
chmod 600 ssl/privkey.pem
chmod 644 ssl/fullchain.pem

# Verify the certificate was created
ls -la ssl/
cat ssl/fullchain.pem
cat ssl/privkey.pem
cat ssl/cert.csr
```

After running these commands, you should see the certificate files in the ssl directory.

7.2.2 Step 2: Modify Nginx Configuration**

Update the existing ragflow.https.conf file to use your IP address.

```
# Go to the ragflow project directory
cd ~/ragflow
pwd
```

```
# Copy the existing HTTPS config as a backup
cp docker/nginx/ragflow.https.conf docker/nginx/ragflow.https.ip.conf

# Edit the configuration file to replace domain with IP and proxy path
sed -i "s/your-ragflow-domain.com/localhost/g" docker/nginx/ragflow.https.ip.conf
sed -i 's|http://ragflow:|http://ragflow-cpu:|' docker/nginx/ragflow.https.ip.conf
# Only if using GPU version, uncomment the following line
#sed -i 's|http://ragflow:|http://ragflow-gpu:|' docker/nginx/ragflow.https.ip.conf

# Verify the changes
cat docker/nginx/ragflow.https.ip.conf
```

7.2.3 Step 3: Update Docker Compose Configuration**

Modify the docker-compose.yml file to mount the self-signed certificates.

```
pwd
cd ~/ragflow/docker

# Add certificate volumes to the ragflow service in docker-compose.yml
# Insert these lines under the existing volumes section for ragflow-cpu and ragflow-gpu
→ services

# For ragflow-cpu service (around line 37-44):
# Add after: - ./entrypoint.sh:/ragflow/entrypoint.sh
    - ./nginx/ssl/fullchain.pem:/etc/nginx/ssl/fullchain.pem:ro
    - ./nginx/ssl/privkey.pem:/etc/nginx/ssl/privkey.pem:ro
    - ./nginx/ragflow.https.ip.conf:/etc/nginx/conf.d/ragflow.conf

# For ragflow-gpu service (around line 86-93):
# Add after: - ./entrypoint.sh:/ragflow/entrypoint.sh
    - ./nginx/ssl/fullchain.pem:/etc/nginx/ssl/fullchain.pem:ro
    - ./nginx/ssl/privkey.pem:/etc/nginx/ssl/privkey.pem:ro
    - ./nginx/ragflow.https.ip.conf:/etc/nginx/conf.d/ragflow.conf

# Verify the docker-compose.yml file syntax
docker compose config
```

7.2.4 Step 4: Restart RAGFlow Services**

Stop the existing services and restart them with the new HTTPS configuration.

```
cd ~/ragflow/docker
# Stop all running services
docker compose down

# Start the services (use --profile cpu or gpu as needed)
```

```
docker compose --profile cpu up -d

# Wait for services to start
sleep 30

# Check that services are running
docker compose ps

# Verify Nginx is listening on ports 80 and 443
docker compose exec ragflow-cpu netstat -tlnp | grep :80
docker compose exec ragflow-cpu netstat -tlnp | grep :443
```

7.2.5 Step 5: Test HTTPS Access**

Verify that the HTTPS setup is working correctly.

```
# Test HTTP to HTTPS redirect
curl -I http://$SERVER_IP

# Test HTTPS connection (ignore certificate warnings)
curl -k -I https://$SERVER_IP

# Check certificate details
openssl s_client -connect $SERVER_IP:443 -servername $SERVER_IP < /dev/null 2>/dev/null |
  openssl x509 -noout -dates -subject

# Verify RagFlow API is accessible over HTTPS
curl -k https://$SERVER_IP/v1/system/healthz
```

7.3 Use Case: Access RagFlow via SSH-Tunnel

This creates an SSH tunnel to access RagFlow securely from a local browser.

```
# Listing: Testing RagFlow via SSH tunnel

ssh -L 9380:localhost:80 EC2-RagFlow-GPU
# Open RagFlow in your browser at http://localhost:9380
exit
```

When the SSH tunnel is established, RagFlow can be accessed via local browser at: <http://localhost:9380>

Result: RagFlow is accessible via SSH tunnel and ready for document processing and RAG operations.

7.4 Summary

RagFlow deployment establishes the core RAG platform with GPU acceleration. The application configures with Docker containerization, HTTPS setup with self-signed certificates, and SSH tunnel access for secure local browser connectivity. The GPU-enabled infrastructure provides enhanced performance for document processing and vector operations.

8 Advanced Features - Document Processing

8.1 Document Ingestion and Processing

This demonstrates uploading and processing documents for RAG functionality with GPU acceleration.

8.1.1 Upload Documents

Access RagFlow via browser and upload sample documents for processing.

8.1.2 Configure Knowledge Base

Create a knowledge base and configure document parsing settings.

8.1.3 Test RAG Queries

Test retrieval-augmented generation with uploaded documents, leveraging GPU acceleration for faster processing.

8.2 Summary

Document processing enables advanced knowledge management while preserving privacy. GPU acceleration provides enhanced performance for embeddings generation and vector similarity searches. Local processing handles sensitive documents without transmitting data to external services.

9 Infrastructure Cleanup

9.1 Cleanup Philosophy

Infrastructure as Code enables rapid deployment, but equally important is the ability to cleanly tear down resources to avoid unnecessary costs. This section provides procedures for complete cleanup of all deployed GPU resources.

9.2 Stop RagFlow Services

Before destroying infrastructure, gracefully stop all running services.

```
# Listing: Stopping RagFlow services
# Gracefully shutting down all containers

ssh EC2-RagFlow-GPU

# Stop RagFlow Docker containers
cd ~/ragflow
docker-compose down

# Verify containers are stopped
docker ps -a

# Optional: Remove Docker images to free disk space
docker system prune -a -f
```

```
# Exit SSH session
exit
```

Result: RagFlow services are stopped and Docker containers are removed from the GPU EC2 instance.

9.3 Terminate GPU EC2 Instance with Terraform

Use Terraform to cleanly destroy all AWS infrastructure resources including the GPU instance.

```
# Listing: Destroying AWS GPU infrastructure
# Removing all Terraform-managed resources

cd "$PROJECT_DIR/terraform-gpu"
pwd
# Review what will be destroyed
terraform plan -destroy -var-file=terraform.tfvars
# Destroy all infrastructure
terraform destroy -var-file=terraform.tfvars
yes

# Verify destruction
terraform show

# List remaining EC2 instances (should show none with EC2-RagFlow-GPU tag)
aws ec2 describe-instances \
--filters "Name>tag:Name,Values=EC2-RagFlow-GPU" \
--query 'Reservations[*].Instances[*].[InstanceId, State.Name]' \
--output text
```

Result: All AWS GPU infrastructure resources are destroyed, including EC2 instance, security groups, and key pairs.

9.4 Clean Up Local SSH Configuration

Remove SSH configurations and keys associated with the destroyed GPU instance.

```
# Listing: Removing local SSH configuration
# Cleaning up SSH keys and config entries

# Remove SSH config entry for EC2-RagFlow-GPU
# Manual edit recommended to preserve other configurations
echo "Removing SSH config entry for EC2-RagFlow-GPU..."
perl -i.bak -ne
↪ 'print unless /^Host EC2-RagFlow-GPU$/. /(^Host / && !/^Host EC2-RagFlow-GPU$/)'
↪ ~./ssh/config

# Verify removal
grep -A 5 "EC2-RagFlow-GPU" ~./ssh/config || echo "SSH config entry removed successfully"

# Optional: Remove SSH keys (only if no longer needed)
# Uncomment the following lines if you want to delete the keys:
# rm -f ~./ssh/ragflow-gpu_key
```

```
# rm -f ~/.ssh/ragflow-gpu_key.pub  
  
# List remaining SSH keys  
ls -la ~/.ssh/ragflow-gpu_key*
```

Result: SSH configuration entries are removed. SSH keys remain available for potential future deployments (remove manually if no longer needed).

9.5 Clean Up Terraform State Files

Clean up Terraform state and temporary files.

```
# Listing: Cleaning up Terraform files  
# Removing state and temporary files  
  
cd "$PROJECT_DIR/terraform-gpu"  
  
# List Terraform state files  
ls -la terraform.tfstate*  
  
# Optional: Back up state files before removal  
mkdir -p ../terraform-gpu-backups  
cp terraform.tfstate* ../terraform-gpu-backups/ 2>/dev/null || echo  
↪ "No state files to backup"  
  
# Remove Terraform state files (only after infrastructure is destroyed)  
rm -f terraform.tfstate  
rm -f terraform.tfstate.backup  
rm -f tfplan*  
rm -f terraform.tfvars  
  
# Clean up Terraform cache  
rm -rf .terraform  
  
# Verify cleanup  
ls -la
```

Result: Terraform state files and temporary files are removed to complete the cleanup process.

9.6 Verify Complete Cleanup

Perform final verification that all resources have been removed.

```
# Listing: Final cleanup verification  
# Confirming all resources are removed  
  
# Check for any remaining EC2 instances  
aws ec2 describe-instances \  
--filters "Name=tag:project,Values=ragflow-gpu" \  
--query  
↪ 'Reservations[*].Instances[*].[InstanceId, State.Name, Tags[?Key==`Name`][0].Value]' \  
--output table
```

```
# Check for remaining security groups
aws ec2 describe-security-groups \
--filters "Name=group-name,Values=RagFlowGPUSecurityGroup" \
--query 'SecurityGroups[*].[GroupId, GroupName]' \
--output table

# Check for remaining key pairs
aws ec2 describe-key-pairs \
--filters "Name=key-name,Values=ragflow-gpu-key" \
--query 'KeyPairs[*].[KeyName, KeyFingerprint]' \
--output table

# Summary
echo =====
echo "Cleanup Verification Complete"
echo =====
echo "If any resources are still listed above,"
echo "they may need manual cleanup via AWS Console"
echo =====
```

Result: Final verification confirms all AWS GPU resources have been properly cleaned up and removed.

9.7 Cost Optimization Notes

Some important considerations for cost management with GPU instances:

- **GPU EC2 Instance:** Charges accrue while the instance is running. GPU instances are significantly more expensive than CPU-only instances. Use ‘terraform destroy’ when not actively using the system.
- **EBS Volumes:** Storage costs continue even when instances are stopped. Destroy volumes when not needed.
- **Elastic IPs:** Unused Elastic IPs incur charges. Ensure they’re released during infrastructure destruction.
- **Data Transfer:** Outbound data transfer from AWS to the internet incurs costs.

For occasional use, consider stopping the instance rather than destroying it:

```
# Stop GPU instance without destroying (preserves data but stops compute charges)
aws ec2 stop-instances --instance-ids "$instance_id"

# Start GPU instance when needed again
aws ec2 start-instances --instance-ids "$instance_id"
```

9.8 Summary

Infrastructure cleanup procedures remove all deployed GPU resources. Terraform destroy commands remove AWS infrastructure, SSH configurations clean from local systems, and verification steps confirm proper resource removal. This prevents unexpected costs after project completion.

This project demonstrates that privacy-focused RAG systems are both technically feasible and practically deployable with GPU acceleration using Ubuntu 24.04 and pre-configured Deep Learning AMIs.

Self-hosted solutions provide alternatives to cloud-only AI services while leveraging GPU performance for enhanced capabilities.

This completes the RagFlow AWS GPU deployment using live-scripting methodology. Each code block is executable with F4 in Emacs using ‘send-line-to-vterm’, or can be copied and pasted into any terminal for the same results.